

NAME ~~~~~ ROLL ~~~~~

This exam has 8 printed page/s. Write your name and roll number on **EVERY SIDE (and not just sheet)**, because we may take apart your answer book and/or xerox it for correction. Write your answer clearly within the spaces provided and on any last blank page. Do not write inside the rectangles to be used for grading. **If you need more space than is provided, you probably made a mistake in interpreting the question.** Start with rough work elsewhere, but you need not attach rough work. Use the marks alongside each question for time management. **Illogical or incoherent answers are worse than wrong answers or even *no* answer, and may fetch negative credit.** You may not use any computing or communication device during the exam. You may use textbooks, class notes written by you, approved material downloaded **prior to the exam** from the course Web page, course news group, or the Internet, or notes made available by me for xeroxing. If you use class notes from other student/s, you must obtain them **prior to the exam** and **write down his/her/their name/s and roll number/s** here.

**1.** We are designing a standard linear chain CRF. Input and prediction are sequences  $\mathbf{x} = (x_1, \dots, x_T)$  and  $\mathbf{y} = (y_1, \dots, y_T) \in \{1, \dots, M\}^T$ .

**1.a** Suppose, from domain knowledge, we know that no valid label sequence  $\mathbf{y}$  can have both states  $a$  and  $b$  occur in it. Label sequences that have neither or one of them are allowed. Design (with explanation) a suitable dynamic programming table  $V(t, m, k)$  for  $k \in [\emptyset, a, b]$  and fill in the expressions to complete the table and indicate how to extract the best predicted sequence  $\mathbf{y}$  subject to this constraint.

		5
--	--	---

$V(t, m, \emptyset)$  records paths that have neither  $a$  nor  $b$  in them.  $V(t, m, a)$  records paths where  $a$  occurs but  $b$  does not.  $V(t, m, b)$  records paths where  $b$  occurs but  $a$  does not. The optimal objective is

$$\max \left\{ \max_m V(T, m, \emptyset), \max_m V(T, m, a), \max_m V(T, m, b) \right\}.$$

We initialize  $V(0, \star, \emptyset) = 0$ ,  $V(0, \star, a) = -\infty$ ,  $V(0, \star, b) = -\infty$ , and, for  $t > 0$ ,

$$V(t, m, \emptyset) = \begin{cases} \max_{m' \in [M]} V(t-1, m', \emptyset) + w \cdot \varphi(x_t, m', m), & m \notin \{a, b\} \\ -\infty, & \text{otherwise} \end{cases}$$

$$V(t, a, a) = \max \begin{cases} \max_{m' \in [M]} V(t-1, m', a) + w \cdot \varphi(x_t, m', a) \\ \max_{m' \in [M]} V(t-1, m', \emptyset) + w \cdot \varphi(x_t, m', a) \end{cases}$$

$$V(t, m, a) = \max_{m' \in [M]} V(t-1, m', a) + w \cdot \varphi(x_t, m', m), \quad m \notin \{a, b\}$$

$$V(t, b, a) = -\infty$$

$$V(t, b, b) = \max \begin{cases} \max_{m' \in [M]} V(t-1, m', b) + w \cdot \varphi(x_t, m', b) \\ \max_{m' \in [M]} V(t-1, m', \emptyset) + w \cdot \varphi(x_t, m', b) \end{cases}$$

$$V(t, m, b) = \max_{m' \in [M]} V(t-1, m', b) + w \cdot \varphi(x_t, m', m), \quad m \notin \{a, b\}$$

$$V(t, a, b) = -\infty$$

- 1.b** Suppose, from domain knowledge, we know that at least one of two designated states  $a$  and  $b$  has to appear at least once in a valid label sequence  $\mathbf{y}$ . Design a suitable dynamic programming table  $V(t, m, k)$  for a suitable index space  $k$  (clearly define it) and fill in the expressions to complete the table and indicate how to extract the best predicted sequence  $\mathbf{y}$  subject to this constraint.

		5
--	--	---

The dynamic programming table has two layers:  $V(t, m, N)$  (seen neither  $a$  nor  $b$  up through time  $t$ ) and  $V(t, m, E)$  (seen either  $a$  or  $b$  or both at least once through time  $t$ ).

$$V(0, \star, N) = 0$$

$$V(\star, a, N) = -\infty$$

$$V(\star, b, N) = -\infty$$

$$V(t, m, N) = \max_{m'} V(t-1, m', N) + w \cdot \varphi(x_t, m', m) \quad m \notin \{a, b\}, t \in [1, T]$$

$$V(0, \star, E) = -\infty$$

$$V(t, m, E) = \max \left\{ \max_{m'} V(t-1, m', N), \max_{m'} V(t-1, m', E) \right\} + w \cdot \varphi(x_t, m', m) \quad m \in \{a, b\}, t \in [1, T]$$

$$V(t, m, E) = \max_{m'} V(t-1, m', E) + w \cdot \varphi(x_t, m', m) \quad m \notin \{a, b\}, t \in [1, T]$$

And finally we pick  $\max_m V(T, m, E)$ .

- 2.** We will develop a fine type formulation based on jointly embedding mention context and type labels, with a few embellishments. Let  $m$  be a mention, whose raw feature vector is collected from text as  $\mathbf{x}_m \in \mathbb{R}^M$ . Let  $\mathbf{U} \in \mathbb{R}^{D \times M}$  project these raw feature vectors to  $\mathbb{R}^D$ .

Suppose there are  $K$  type labels. Type  $k$  is represented by a vector  $\mathbf{y}_k \in \{0, 1\}^K$  where  $y_k(k) = 1$  and all  $y_k(\neq k) = 0$ . Type vectors  $\mathbf{y}$  are projected by matrix  $\mathbf{V} \in \mathbb{R}^{D \times K}$  into  $D$ -dimensional space.

- 2.a** The score of type  $k$ , given mention  $m$ , is modeled as the dot product between the two projections. Write down the score and call it  $f_m(k)$ .

		1
--	--	---

$$f_m(k) = (\mathbf{U}\mathbf{x}_m) \cdot (\mathbf{V}\mathbf{y}_k)$$

- 2.b** Given training instance  $(m_i, k_i)$ , the hinge loss would usually be defined as  $\ell_i(k) = \max\{0, 1 + f_{m_i}(k) - f_{m_i}(k_i)\}$ . But typical type systems have redundant and overlapping types. If  $k$  is “very similar” to  $k_i$  in terms of the entities they contain, the above hinge loss may be unfair. Suppose  $E_k$  is the set of entities contained in type  $k$ . Suggest and justify a margin  $\Delta_k(k')$  that addresses this problem. There may be many acceptable solutions.

		2
--	--	---

Here is one proposal:

$$w_{kk'} = \frac{1}{2} \left( \frac{|E_k \cap E_{k'}|}{|E_k|} + \frac{|E_k \cap E_{k'}|}{|E_{k'}|} \right)$$

$$\Delta_k(k') = \frac{1}{\spadesuit + w_{kk'}}$$

where  $\spadesuit > 0$  is some tuned constant. If  $E_k$  and  $E_{k'}$  overlap a lot,  $\Delta_k(k')$  should be small, and vice versa.

- 2.c** In reality more than one type labels may be active at a mention. Suppose training data comes in the form of certified present and absent types  $(m_i, K_i^+, K_i^-)$ . Complete the following loss function for the  $i$ th instance:

$$\ell_i = \sum_{k \in K_i^+} \sum_{k' \in K_i^-} \underbrace{\hspace{10em}}_{\clubsuit}.$$

		1
--	--	---

$$\ell_i = \sum_{k \in K_i^+} \sum_{k' \in K_i^-} \max\{0, \Delta_k(k') + f_{m_i}(k') - f_{m_i}(k)\} \quad (1)$$

- 2.d** Express the rank  $R_{m_i}(k^*)$  of a correct type label  $k^*$  as a (possibly discontinuous) function of all the label scores  $f_{m_i}(k) : k = 1, \dots, K$ . The label with largest score has rank 0. The loss can then be made rank-sensitive by writing

$$\ell_i = \sum_{k \in K_i^+} \sum_{k' \in K_i^-} \clubsuit R_{m_i}(k).$$

Approximate the rank using a smooth function that can facilitate learning by gradient descent.

		1
--	--	---

Actually, for every mention and every correct label  $k^*$ , we should be counting how many *incorrect* labels defeat it in terms of scores. We should not be comparing two correct labels. Let the sets of correct and incorrect labels be  $K_i^+, K_i^-$ .

$$R_{m_i}(k^*) = \sum_{k' \in K_i^-} \begin{cases} 1, & \Delta_{k^*}(k') + f_{m_i}(k') > f_{m_i}(k^*), \\ 0, & \text{otherwise} \end{cases}$$

The smooth version can be obtained via a sigmoid:

$$R_{m_i}(k^*) = \sum_{k' \in K_i^-} \sigma(\blacklozenge(\Delta_{k^*}(k') + f_{m_i}(k') - f_{m_i}(k^*)) + \heartsuit),$$

where  $\blacklozenge > 0$  and  $\heartsuit$  are hyperparameters. These sums-of-sigmoids are generally hard to train, though.

- 2.e** We need to be robust to noisy training data. If a corpus is annotated with mentions  $m$  of entities  $e$ , and from the KG we know a set of types  $K_e$  to which  $e$  belongs, we should not use  $K_e$  as  $K^+$  for all mentions of  $e$ . To tackle this problem, we divide mention instances into two kinds: *clean* mentions  $\mathcal{M}_c$  where the KG connects  $e$  to exactly one path in the type hierarchy, and *noisy* mentions where more than one paths are connected to the entity, but only some of them may be active in a specific mention. We will write the overall loss objective as

$$\min_{\mathbf{U}, \mathbf{V}} \left[ \sum_{(m_i, K_i^+, K_i^-) \in \mathcal{M}_c} \boxed{\text{clean mention loss}} + \sum_{(m_i, K_i^+, K_i^-) \in \mathcal{M}_n} \boxed{\text{noisy mention loss}} \right]$$

For the clean mentions we will use the above loss function. For noisy mentions, we will insist only that the *largest* score from among  $K_i^+$  should beat all scores from  $K_i^-$ . Write out in detail the noisy loss function.

		2
--	--	---

We need to change the symmetric sum-sum form in (1) to a sum-max form: For each bad type, *some* good type should beat it.

$$\ell_i = \sum_{k' \in K_i^-} \max \left\{ 0, f_{m_i}(k') - \max_{k^* \in K_i^+} [f_{m_i}(k^*) - \Delta_{k^*}(k')] \right\}$$

This is not the only reasonable proposal.

- 3.** We are given a document with  $N$  entity mentions with contexts. Our goal is to infer entity labels  $\mathbf{y} = y_1, \dots, y_N$ . In the star model, we choose

$$\hat{y}_i = \operatorname{argmax}_{y_i} \left[ \phi_i(y_i) + \sum_{j \neq i} \max_{y_j} \psi_{ij}(y_i, y_j) \right].$$

**3.a** Complete the following pseudocode to compute  $\hat{\mathbf{y}}$ .

```

for each  $i = 1, \dots, N$  do
   $bestiLabel \leftarrow \text{NULL}$ ,  $bestiScore \leftarrow -\infty$ 
  for each possible label  $y_i$  do
     $yiScore \leftarrow \text{~~~~~}$ 
    for  $j = 1, \dots, N, j \neq i$  do
       $bestjiSupport \leftarrow \text{~~~~~}$ 
      for each possible label  $y_j$  do
        if  $bestjiSupport < \text{~~~~~}$  then
           $bestjiSupport \leftarrow \text{~~~~~}$ 
         $yiScore \leftarrow yiScore + \text{~~~~~}$ 
      if  $bestiScore < \text{~~~~~}$  then
         $bestiScore \leftarrow \text{~~~~~}$ 
         $bestiLabel \leftarrow \text{~~~~~}$ 
    set  $\hat{y}_i \leftarrow bestiLabel$ 
return  $\hat{\mathbf{y}}$ 

```

		4
--	--	---

```

for each  $i = 1, \dots, N$  do
   $bestiLabel \leftarrow \text{NULL}$ ,  $bestiScore \leftarrow -\infty$ 
  for each possible label  $y_i$  do
     $yiScore \leftarrow \phi_i(y_i)$ 
    for  $j = 1, \dots, N, j \neq i$  do
       $bestjiSupport \leftarrow -\infty$ 
      for each possible label  $y_j$  do
        if  $bestjiSupport < \psi_{ij}(y_i, y_j)$  then
           $bestjiSupport \leftarrow \psi_{ij}(y_i, y_j)$ 
         $yiScore \leftarrow yiScore + \text{~~~~~}$ 
      if  $bestiScore < \text{~~~~~}$  then
         $bestiScore \leftarrow \text{~~~~~}$ 
         $bestiLabel \leftarrow y_i$ 
    set  $\hat{y}_i \leftarrow bestiLabel$ 
return  $\hat{\mathbf{y}}$ 

```

**3.b** A potential problem with this formulation is that there is no guarantee that  $\hat{y}_j$  will be the best supporting label for  $\hat{y}_i$ . We will solve  $N$  separate problems, centered on mentions  $i = 1, \dots, N$ , to infer  $\mathbf{y}^{(i)} = (y_1^{(i)}, \dots, y_N^{(i)})$ , and try to make these solutions consistent with a single global solution  $\mathbf{y}^{(0)}$ . For simplicity assume each  $y_i \in \{1, \dots, M\}$ . We will change the signature of  $y_i$  from an integer in  $[1, M]$  to a 1-hot vector  $Y_i \in \{0, 1\}^M$ . Let  $\phi_{im} = \phi_i(m)$  be the local potential of node  $i$ . What is the local score at node  $i$  as a function of  $\phi_i \in \mathbb{R}^M$  and  $Y_i$ ?

		1
--	--	---

The local score (log node potential) at mention (node)  $i$  is  $\phi_i \cdot Y_i$ .

- 3.c** If  $\psi_{ij}$  is represented by a real  $M \times M$  matrix, write down the potential for edge  $(i, j)$  with the nodes having 1-hot labels  $Y_i$  and  $Y_j$ .

		2
--	--	---

The log edge potential between nodes  $i$  and  $j$  is  $Y_i^\top \psi_{ij} Y_j$ .

- 3.d** Relax  $Y_i$  from 1-hot to the unit simplex  $\Delta_M$  in  $M$  dimensions, i.e.,  $Y_{im} \in \mathbb{R}$ ,  $Y_{im} \geq 0$  and  $\sum_m Y_{im} = 1$ . Complete the objective for the local optimization for the  $i$ th problem:

$$\max_{Y_i \in \Delta_M} \left[ \phi_i \cdot \text{~~~~~} + \sum_{j \neq i} \max_{\text{~~~~~} \in \Delta_M} \text{~~~~~} \right]$$

What solver can you use to maximize this objective? (Extra credit: Is the optimization convex? If not, can you make it convex by introducing additional variables and constraints?)

		2
--	--	---

$$\max_{Y_i \in \Delta_M} \left[ \phi_i \cdot Y_i + \sum_{j \neq i} \max_{Y_j \in \Delta_M} Y_i^\top \psi_{ij} Y_j \right]$$

The optimization is convex if all feasible regions are convex (which holds) and the objective to maximize is concave (which does not hold in general because of the  $Y_i^\top \psi_{ij} Y_j$  terms). The standard technique to turn the slave optimizations convex (in fact, linear) is to introduce auxiliary variables  $Z_{i,m;i',m'} \in \{0, 1\}$  (or in  $[0, 1]$  when relaxed) such that  $Y_{i,m} = \sum_{i',m'} Z_{i,m;i',m'}$  etc., and then write the objective as a linear function of  $Y$  and  $Z$ . This results in additional (in)equality constraints coupling  $Y, Z$  but the whole slave optimization can be solved by an LP solver.

- 3.e** Following through with our plan above, we will solve  $N$  problems. In the  $i$ th problem, the relaxed label variables will be  $\mathbf{Y}^{(i)} = (Y_1^{(i)}, \dots, Y_N^{(i)}) \in \Delta_M^N$ . We will then create a penalty if these solutions deviate from a global solution  $\mathbf{Y}^{(0)} = (Y_1^{(0)}, \dots, Y_N^{(0)}) \in \Delta_M^N$ . Complete the following global objective:

$$\min_{\lambda \in \Delta_N} \max_{\substack{\mathbf{Y}^{(0)} \in \Delta_M^N \\ \{\mathbf{Y}^{(i)} \in \Delta_M^N : i=1, \dots, N\}}} \sum_{i=1}^N \left( \text{~~~~~} + \sum_{j \neq i} \text{~~~~~} \right) - \sum_{i=1}^N \lambda_i \|\mathbf{Y}^{(i)} - \text{~~~~~}\|_1.$$

Here  $\|C\|_1$  is the L1 norm of  $C$ . (This may not be the way dual decomposition is usually set up. If you prefer you may set up the global objective in a different way, as long as it satisfies our goals.)

		1
--	--	---

$$\min_{\boldsymbol{\lambda} \in \Delta_N} \max_{\substack{\mathbf{Y}^{(0)} \in \Delta_M^N \\ \{\mathbf{Y}^{(i)} \in \Delta_M^N : i=1, \dots, N\}}} \sum_{i=1}^N \left( \phi_i \cdot \underbrace{Y_i^{(i)}} + \sum_{j \neq i} \max_{\underbrace{Y_j^{(i)} \in \Delta_M}} Y_i^{(i)\top} \psi_{ij} Y_j^{(i)} \right) - \clubsuit \sum_{i=1}^N \lambda_i \|\mathbf{Y}^{(i)} - \underbrace{\mathbf{Y}^{(0)}}\|_1.$$

Note that the penalty has to be *subtracted*, and it might help to multiply by some balancing factor  $\clubsuit$ .

**3.f** Given the high-level pseudocode below:

```

choose initial  $\boldsymbol{\lambda}, \mathbf{Y}^{(0)}$ 
for some number of iterations do
  for each problem indexed  $i$  do
    fix current  $\boldsymbol{\lambda}, \mathbf{Y}^{(0)}$  and solve for next  $\mathbf{Y}^{(i)}$ 
  update  $\mathbf{Y}^{(0)}$ 
  update  $\boldsymbol{\lambda}$ 

```

write down how you would use standard optimizers (or simple update expressions) to solve the three key steps.

		3
--	--	---

**Solving for  $\mathbf{Y}^{(i)}$ :** In the pseudocode for choosing  $\hat{\mathbf{y}}$ , we could search through all possible  $y_i$ , and let the current choice of  $y_i$  force the best possible discrete  $y_j$ s. It is not necessarily optimal to follow a similar recipe, but it may behave reasonably in practice:

- First optimize  $\operatorname{argmax}_{Y_i^{(i)} \in \Delta_M} \phi_i Y_i^{(i)} - \lambda_i \|Y_i^{(i)} - Y_i^{(0)}\|_1$  using an LP solver. Note that  $\lambda_i$  and  $Y_i^{(0)}$  are frozen at the moment.
- Now for each  $j \neq i$ , choose  $\operatorname{argmax}_{Y_j^{(i)} \in \Delta_M} \boxed{Y_i^{(i)\top} \psi_{ij}} Y_j^{(i)}$ . Now that  $Y_i^{(i)}$  is also fixed momentarily, the boxed part is a constant matrix. So this problem is also easily solved via LP.

Projected gradient descent may also work.

**Updating  $\mathbf{Y}^{(0)}$ :** This is just finding a weighted medoid, easily solved via LP.

**Updating  $\boldsymbol{\lambda}$ :** At this point we have constants  $a_i = \|\mathbf{Y}^{(i)} - \mathbf{Y}^{(0)}\|_1 \geq 0$ , and must find  $\operatorname{argmin}_{\boldsymbol{\lambda} \in \Delta_N} \sum_i a_i \lambda_i$ . Naturally we would set  $\lambda_{i^*} = 1$  for  $i^* = \operatorname{argmin}_i a_i$  and  $\lambda_i = 0$  for all  $i \neq i^*$ . In other words, every outer iteration,  $\boldsymbol{\lambda}$  would get set to a 1-hot vector, which may not be a good thing. Encouraging some entropy to  $\boldsymbol{\lambda}$  may help stability in practice.

Dual decomposition differs from the above in various ways. We would explicitly introduce decision variables  $Z_{ij}$  for every edge, leading to a linear objective  $\psi_{ij} \odot Z_{ij}$ .  $Z$  and  $Y$  would be coupled by marginal inequalities. Also, we would use  $\boldsymbol{\lambda} \in \mathbb{R}^{N \times N \times M}$ , not take norm errors, but instead write something like

Roll:

Name:

---

$\sum_{n=1}^N \sum_{i=1}^N \sum_{j=1}^M \lambda_{nij} (Y_{ij}^{(n)} - Y_{ij}^{(0)})$ . But this would let us solve all local problems exactly via LP.

**Total: 30**